

Рубежный контроль No2

Тема: Методы построения моделей машинного обучения

Федоров И.Н. ИУ5-61Б

Загрузка необходимых библиотек:

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
from sklearn import preprocessing
from sklearn import svm
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, cla
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, mea
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.linear_model import LogisticRegression, LogisticRegressionC
from sklearn.ensemble import RandomForestClassifier

from sklearn.ensemble import AdaBoostClassifier
from sklearn.impute import SimpleImputer
```

```
In [2]: data = pd.read_csv('./HRDataset_v14.csv', sep=",")
TARGET_COL_NAME = 'RecruitmentSource'
TARGET_IS_NUMERIC = data[TARGET_COL_NAME].dtype != 'O'
TARGET_IS_NUMERIC
```

Out[2]: False

```
In [3]: data
```

```
Out[3]:
```

	Employee_Name	EmpID	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	
0	Adinolfi, Wilson K	10026	0	0	1	1	5	
1	Ait Sidi, Karthikeyan	10084	1	1	1	5	3	
2	Akinkuolie, Sarah	10196	1	1	0	5	5	
3	Alagbe, Trina	10088	1	1	0	1	5	
4	Anderson, Carol	10069	0	2	0	5	5	
...	
306	Woodson, Jason	10135	0	0	1	1	5	

307	Ybarra, Catherine	10301	0	0	0	5	5
308	Zamora, Jennifer	10010	0	0	0	1	3
309	Zhou, Julia	10043	0	0	0	1	3
310	Zima, Colleen	10271	0	4	0	1	5

311 rows × 36 columns

In [4]: `data.shape`

Out[4]: (311, 36)

In [5]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Employee_Name                        311 non-null    object
1   EmpID                               311 non-null    int64
2   MarriedID                           311 non-null    int64
3   MaritalStatusID                     311 non-null    int64
4   GenderID                            311 non-null    int64
5   EmpStatusID                         311 non-null    int64
6   DeptID                              311 non-null    int64
7   PerfScoreID                         311 non-null    int64
8   FromDiversityJobFairID              311 non-null    int64
9   Salary                              311 non-null    int64
10  Termd                               311 non-null    int64
11  PositionID                          311 non-null    int64
12  Position                             311 non-null    object
13  State                               311 non-null    object
14  Zip                                  311 non-null    int64
15  DOB                                 311 non-null    object
16  Sex                                  311 non-null    object
17  MaritalDesc                         311 non-null    object
18  CitizenDesc                         311 non-null    object
19  HispanicLatino                     311 non-null    object
20  RaceDesc                           311 non-null    object
21  DateofHire                         311 non-null    object
22  DateofTermination                  104 non-null    object
23  TermReason                         311 non-null    object
24  EmploymentStatus                   311 non-null    object
25  Department                         311 non-null    object
26  ManagerName                        311 non-null    object
27  ManagerID                          303 non-null    float64
28  RecruitmentSource                   311 non-null    object
29  PerformanceScore                    311 non-null    object
30  EngagementSurvey                    311 non-null    float64
31  EmpSatisfaction                     311 non-null    int64
32  SpecialProjectsCount                311 non-null    int64
33  LastPerformanceReview_Date          311 non-null    object
34  DaysLateLast30                      311 non-null    int64
35  Absences                            311 non-null    int64
dtypes: float64(2), int64(16), object(18)
memory usage: 87.6+ KB
```

In [6]:


```
# проверим есть ли пропущенные значения
data.isnull().sum()
```

```
Out[6]: Employee_Name      0
EmpID      0
MarriedID   0
MaritalStatusID  0
GenderID    0
EmpStatusID 0
DeptID      0
PerfScoreID 0
FromDiversityJobFairID 0
Salary      0
Termd       0
PositionID  0
Position    0
State       0
Zip         0
DOB         0
Sex         0
MaritalDesc 0
CitizenDesc 0
HispanicLatino 0
RaceDesc    0
DateofHire   0
DateofTermination 207
TermReason   0
EmploymentStatus 0
Department   0
ManagerName  0
ManagerID    8
RecruitmentSource 0
PerformanceScore 0
EngagementSurvey 0
EmpSatisfaction 0
SpecialProjectsCount 0
LastPerformanceReview_Date 0
DaysLateLast30 0
Absences     0
dtype: int64
```

Удалим колонки, которые не влияют на целевой признак

Построим heatmap для лучшего визуального представления всех корреляций

```
In [7]: cmap = sns.cm.rocket_r
ax = sns.heatmap(data.corr(), cmap=cmap)
```



```
In [8]: data = data.drop(columns=['Employee_Name', 'EmpID', 'DateofTermination'],
data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MarriedID                            311 non-null    int64
1   MaritalStatusID                     311 non-null    int64
2   GenderID                            311 non-null    int64
3   EmpStatusID                         311 non-null    int64
4   DeptID                             311 non-null    int64
5   PerfScoreID                        311 non-null    int64
6   FromDiversityJobFairID              311 non-null    int64
7   Salary                             311 non-null    int64
8   Termd                              311 non-null    int64
9   PositionID                          311 non-null    int64
10  Position                             311 non-null    object
11  State                               311 non-null    object
12  Zip                                 311 non-null    int64
13  DOB                                311 non-null    object
14  Sex                                 311 non-null    object
15  MaritalDesc                         311 non-null    object
16  CitizenDesc                         311 non-null    object
17  HispanicLatino                      311 non-null    object
18  RaceDesc                           311 non-null    object
19  DateofHire                         311 non-null    object
20  TermReason                         311 non-null    object
21  EmploymentStatus                   311 non-null    object
22  Department                         311 non-null    object
23  ManagerName                        311 non-null    object
24  RecruitmentSource                  311 non-null    object
25  PerformanceScore                   311 non-null    object
26  EngagementSurvey                   311 non-null    float64
27  EmpSatisfaction                    311 non-null    int64
28  SpecialProjectsCount               311 non-null    int64
29  LastPerformanceReview_Date         311 non-null    object
30  DaysLateLast30                     311 non-null    int64
```

```
31 Absences 311 non-null int64
dtypes: float64(1), int64(15), object(16)
memory usage: 77.9+ KB
```

Обработка пропусков

In [9]:

```
# Импультация наиболее частыми значениями
imp = SimpleImputer(missing_values=np.nan, strategy='most_frequent')

imputed = {}

for col in data:
    contains_nan = data[col].isnull().sum() != 0
    if contains_nan:
        data_imp = data[[col]]
        data_imp = imp.fit_transform(data_imp)
        imputed[col] = data_imp

for col_name in imputed:
    df = pd.DataFrame({col_name:imputed[col_name].T[0]})
    data[col_name] = df.copy()

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 311 entries, 0 to 310
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	MarriedID	311 non-null	int64
1	MaritalStatusID	311 non-null	int64
2	GenderID	311 non-null	int64
3	EmpStatusID	311 non-null	int64
4	DeptID	311 non-null	int64
5	PerfScoreID	311 non-null	int64
6	FromDiversityJobFairID	311 non-null	int64
7	Salary	311 non-null	int64
8	TermId	311 non-null	int64
9	PositionID	311 non-null	int64
10	Position	311 non-null	object
11	State	311 non-null	object
12	Zip	311 non-null	int64
13	DOB	311 non-null	object
14	Sex	311 non-null	object
15	MaritalDesc	311 non-null	object
16	CitizenDesc	311 non-null	object
17	HispanicLatino	311 non-null	object
18	RaceDesc	311 non-null	object
19	DateofHire	311 non-null	object
20	TermReason	311 non-null	object
21	EmploymentStatus	311 non-null	object
22	Department	311 non-null	object
23	ManagerName	311 non-null	object
24	RecruitmentSource	311 non-null	object
25	PerformanceScore	311 non-null	object
26	EngagementSurvey	311 non-null	float64
27	EmpSatisfaction	311 non-null	int64
28	SpecialProjectsCount	311 non-null	int64
29	LastPerformanceReview_Date	311 non-null	object
30	DaysLateLast30	311 non-null	int64
31	Absences	311 non-null	int64

```
dtypes: float64(1), int64(15), object(16)
memory usage: 77.9+ KB
```

Кодирование строковых признаков (LabelEncoding)

```
In [10]: not_number_cols = data.select_dtypes(include=['object'])
         number_cols = data.select_dtypes(exclude=['object'])
```

```
In [11]: le = preprocessing.LabelEncoder()

         for col_name in not_number_cols:
             data[col_name] = le.fit_transform(data[col_name])

         data
```

```
Out[11]:
```

	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversity
0	0	0	1	1	5	4	
1	1	1	1	5	3	3	
2	1	1	0	5	5	3	
3	1	1	0	1	5	3	
4	0	2	0	5	5	3	
...	
306	0	0	1	1	5	3	
307	0	0	0	5	5	1	
308	0	0	0	1	3	4	
309	0	0	0	1	3	3	
310	0	4	0	1	5	3	

311 rows × 32 columns

Масштабируем числовые данные

```
In [12]: scaler = preprocessing.MinMaxScaler()

         number_fields_source = number_cols.loc[:, number_cols.columns != TARGET_COLUMN]

         for col_name in number_fields_source:
             data[col_name] = scaler.fit_transform(data[[col_name]])

         data
```

```
Out[12]:
```

	MarriedID	MaritalStatusID	GenderID	EmpStatusID	DeptID	PerfScoreID	FromDiversity
0	0.0	0.00	1.0	0.0	0.8	1.000000	
1	1.0	0.25	1.0	1.0	0.4	0.666667	
2	1.0	0.25	0.0	1.0	0.8	0.666667	
3	1.0	0.25	0.0	0.0	0.8	0.666667	
4	0.0	0.50	0.0	1.0	0.8	0.666667	

...
306	0.0	0.00	1.0	0.0	0.8	0.666667
307	0.0	0.00	0.0	1.0	0.8	0.000000
308	0.0	0.00	0.0	0.0	0.4	1.000000
309	0.0	0.00	0.0	0.0	0.4	0.666667
310	0.0	1.00	0.0	0.0	0.8	0.666667

311 rows × 32 columns

Делим выборку на обучающую и тестовую

```
In [13]: target = data[TARGET_COL_NAME]
data_X_train, data_X_test, data_y_train, data_y_test = train_test_split(
    data, target, test_size=0.2, random_state=1)
```

```
In [14]: data_X_train.shape, data_y_train.shape
```

```
Out[14]: ((248, 32), (248,))
```

```
In [15]: data_X_test.shape, data_y_test.shape
```

```
Out[15]: ((63, 32), (63,))
```

```
In [16]: np.unique(target)
```

```
Out[16]: array([0, 1, 2, 3, 4, 5, 6, 7, 8])
```

Логистическая регрессия

```
In [17]: svr_1 = LogisticRegression(solver='lbfgs', max_iter=1000)
svr_1.fit(data_X_train, data_y_train)
```

C:\Users\pstri\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
Out[17]: LogisticRegression(max_iter=1000)
```

```
In [18]: data_y_pred_1 = svr_1.predict(data_X_test)
accuracy_score(data_y_test, data_y_pred_1)
```

Out[18]: 0.7619047619047619

```
In [19]: f1_score(data_y_test, data_y_pred_1, average='micro')
```

Out[19]: 0.7619047619047619

```
In [20]: f1_score(data_y_test, data_y_pred_1, average='macro')
```

Out[20]: 0.7208312792201521

```
In [21]: f1_score(data_y_test, data_y_pred_1, average='weighted')
```

Out[21]: 0.7563589699202566

```
In [22]: svr_2 = LogisticRegression(solver='lbfgs', max_iter=10000)
svr_2.fit(data_X_train, data_y_train)
```

C:\Users\pstri\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Out[22]: LogisticRegression(max_iter=10000)

```
In [23]: data_y_pred_2 = svr_2.predict(data_X_test)
accuracy_score(data_y_test, data_y_pred_2)
```

Out[23]: 0.8095238095238095

```
In [24]: f1_score(data_y_test, data_y_pred_2, average='micro')
```

Out[24]: 0.8095238095238095

```
In [25]: f1_score(data_y_test, data_y_pred_2, average='macro')
```

Out[25]: 0.5976659982174688

```
In [26]: f1_score(data_y_test, data_y_pred_2, average='weighted')
```

Out[26]: 0.7990160710748947

Случайный лес


```
In [27]: RT = RandomForestClassifier(n_estimators=15, random_state=123)
RT.fit(data_X_train, data_y_train)
```

```
Out[27]: RandomForestClassifier(n_estimators=15, random_state=123)
```

```
In [28]: accuracy_score(data_y_test, RT.predict(data_X_test))
```

```
Out[28]: 0.7777777777777778
```

```
In [29]: f1_score(data_y_test, data_y_pred_1, average='micro')
```

```
Out[29]: 0.7619047619047619
```

```
In [30]: f1_score(data_y_test, data_y_pred_1, average='macro')
```

```
Out[30]: 0.7208312792201521
```

```
In [31]: f1_score(data_y_test, data_y_pred_1, average='weighted')
```

```
Out[31]: 0.7563589699202566
```

```
In [32]: RT = RandomForestClassifier(n_estimators=30, random_state=123)
RT.fit(data_X_train, data_y_train)
```

```
Out[32]: RandomForestClassifier(n_estimators=30, random_state=123)
```

```
In [33]: accuracy_score(data_y_test, RT.predict(data_X_test))
```

```
Out[33]: 0.873015873015873
```

```
In [34]: f1_score(data_y_test, data_y_pred_1, average='micro')
```

```
Out[34]: 0.7619047619047619
```

```
In [35]: f1_score(data_y_test, data_y_pred_1, average='macro')
```

```
Out[35]: 0.7208312792201521
```

```
In [36]: f1_score(data_y_test, data_y_pred_1, average='weighted')
```

```
Out[36]: 0.7563589699202566
```

Выводы

При использовании логистической регрессии наилучшую точность (0.809) показала модель с параметром `max_iter=10000`. При использовании метода "Случайный

лес" получилось добиться более высокого показателя точности (0.873), поэтому в целом предпочтительнее использовать его.